



# Gamefest

MICROSOFT GAME TECHNOLOGY CONFERENCE 2 0 0 8

Microsoft

# Practical Parallel Rendering with DirectX 9 and 10

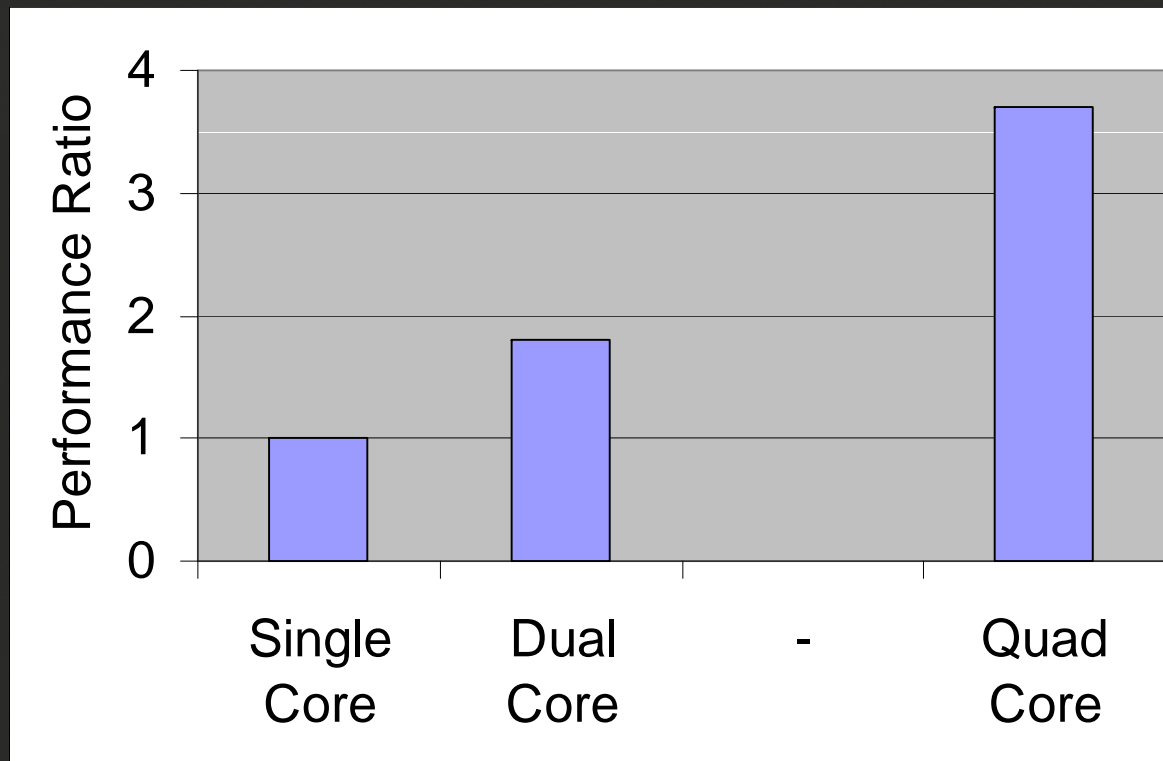
## Windows PC Command Buffers

Vincent Scheib  
Architect  
Gamebryo  
Emergent Game Technologies



# Introduction

- Take advantage of multiple cores with parallel rendering
- Performance should scale by number of cores



Observed data  
from this project,  
details follow

# Presentation Outline

Motivation and problem definition

- Command buffers
  - Requirements
  - Implementation
  - Handling effects and resources
- Application models
- Integrating to existing code
- Prototype results
- Future work

# Motivation

- Take advantage of multi core machines
  - 40% machines have 2+ physical CPUs (steamJul08)
- Rendering can have high CPU cost
- Direct3D 11 display lists coming, but want support for Direct3D 9 and 10 now
  - Currently 81% DX9 HW, 9% DX10 HW (steamJul08)
  - Rough DX9 HW forecast: 2011 ~30% (emergent)
  - Asia HW trends lag somewhat



# Multithreaded DX Device?

- DirectX 9 and 10 primarily designed for single-threaded game architectures
- Multithreaded mode incurs overhead
  - Cuts FPS roughly in half on DX9 for a CPU render call bound application
- DX is Stateful
  - Requires additional synchronization for parallel rendering

# Ideal Scenario

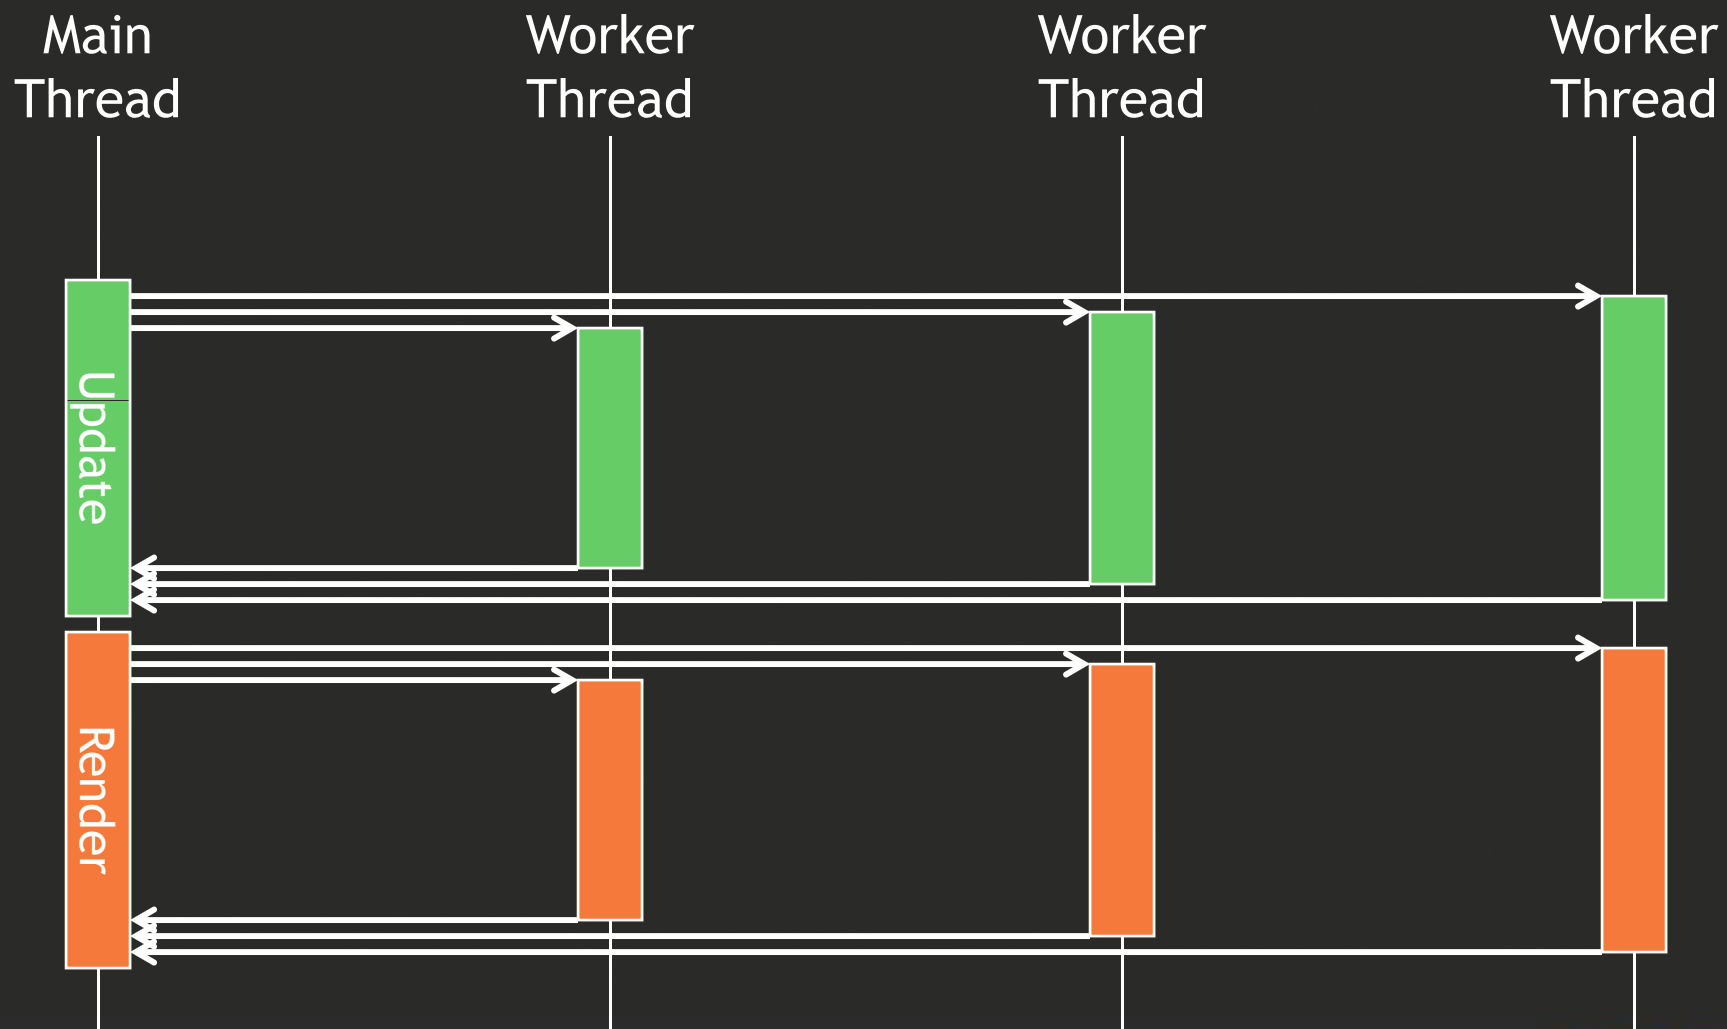
- One thread per hardware thread
- Application manages dispatching work to multiple threads
- Rendering data completely prepared, ready to be sent to single-threaded D3D device
  - Function calls, conditionals, and final matrix multiplies are wasted time on a D3D device thread

# Reality

- Update()
  - Seldomly generates coherent data in API specific format.
- Render()
  - Still works between calls to DX API



# Going Wide



# Command Buffers

- Record calls to D3D
  - Store in a command buffer
  - Can be done concurrently on multiple threads, to multiple command buffers
- Playback D3D commands
  - Efficiently on main thread
  - Exact data for DX API
  - Coherent in memory
- Clean and modular point to integrate to application

# Command Buffer Requirements

- Minimal modifications to rendering code
  - Most code uses pointer to D3DDevice
  - Parameters from stack, e.g., D3DRECT
  - Support most of the device API
    - Draw calls, setting state, constants, shaders, textures, stream source, and so on
  - Support effects
- Playback does not modify buffer
- Playback is ideal performance

# Command Buffer Allowances

- No support for:
  - Create methods
  - Get methods
  - Miscellaneous other functions that return values
    - QueryInterface, ShowCursor

# Command Buffer: Nice to Have

- Buffers played back multiple times
- Optimization of buffers
  - Remove redundant state calls
    - Offload main thread by doing this on recorder threads
  - Reordering of sort independent draw calls



# Design: Recording

- Wrap every API call
  - Unsupported calls, return error
  - Supported calls
    - Store enumeration for call into buffer
    - Store parameters into buffer
    - Make copies of non-reference counted objects such as D3DMATRIX, D3DRECT, shader constants, and so on

# Design: Playback

- Playback, read from buffer, and
  - select function call pointer from table given token
  - each playback function unpacks parameters buffer

# Recording Example

```
virtual HRESULT STDMETHODCALLTYPE DrawPrimitive(  
    D3DPRIMITIVETYPE PrimitiveType,  
    UINT StartVertex,  
    UINT PrimitiveCount)  
{  
    m_pCommandBuffer->Put(CBD3D_COMMANDS::DrawPrimitive);  
  
    m_pCommandBuffer->Put(PrimitiveType);  
    m_pCommandBuffer->Put(StartVertex);  
    m_pCommandBuffer->Put(PrimitiveCount);  
    return D3D_OK;  
}
```

# Playback Example

```
void CBPlayer9::DoDrawPrimitive()  
{  
    D3DPRIMITIVETYPE arg1;  
    m_pCommandBuffer->Get(&arg1);  
  
    UINT arg2;  
    m_pCommandBuffer->Get(&arg2);  
  
    UINT arg3;  
    m_pCommandBuffer->Get(&arg3);  
    if(FAILED(m_pDevice->DrawPrimitive(arg1, arg2, arg3)))  
        OutputDebugStringA(__FUNCTION__ " failed in playback\n");  
}
```



# Effects: Problem

- Effect takes pointer to device at creation
- Effect then creates resources
- At render, effect should use our recorder
- Our recording device cannot create resources



# Effects: Solutions

1. Create FX with command buffer device
  - Fails: needs real device for initialization
2. Wrap and record FX calls and play them back
  - Inefficient
3. Give FX **EffectStateManager** class to redirect calls to command buffer, give it real device for initialization
  - Disables FX use of state blocks
4. Create **redirecting** device
  - Acts as real device at init, command buffer device at render time

# Resource Management

- Multiple threads wish to:
  - Create resources (e.g., background loading)
  - Update resources (e.g., dynamic geometry)
- App must use playback thread only to modify resources
  - App specific logic
    - Deferred creation, double buffering
  - Support in command buffers (next slide)

# Resource Management (2)

- Command buffer library could encapsulate details
  - (This is **Future Work**)
- Gamebryo Volatile Type Buffers
  - D3DUSAGE: WRITEONLY | DYNAMIC  
D3DLOCK: NOOVERWRITE, DISCARD
  - Lock() is stored into command buffer
  - Memory allocated from command buffer, returned from Lock()
  - At playback, true lock is performed
- Gamebryo Mutable Type Buffers:
  - CPU read and infrequent access
  - Backing store required, copied on each Lock()

# Implementation Considerations

- Ease of changing implementation
  - Macros provide implementation
  - Preprocessor & Beautifier produce debuggable code
  - Many macro permutations required (~40) for different argument count and return type
    - Generated from Excel
- Function overloading to store non ref counted parameters
  - Everything but shader constants then stored with same function signature.

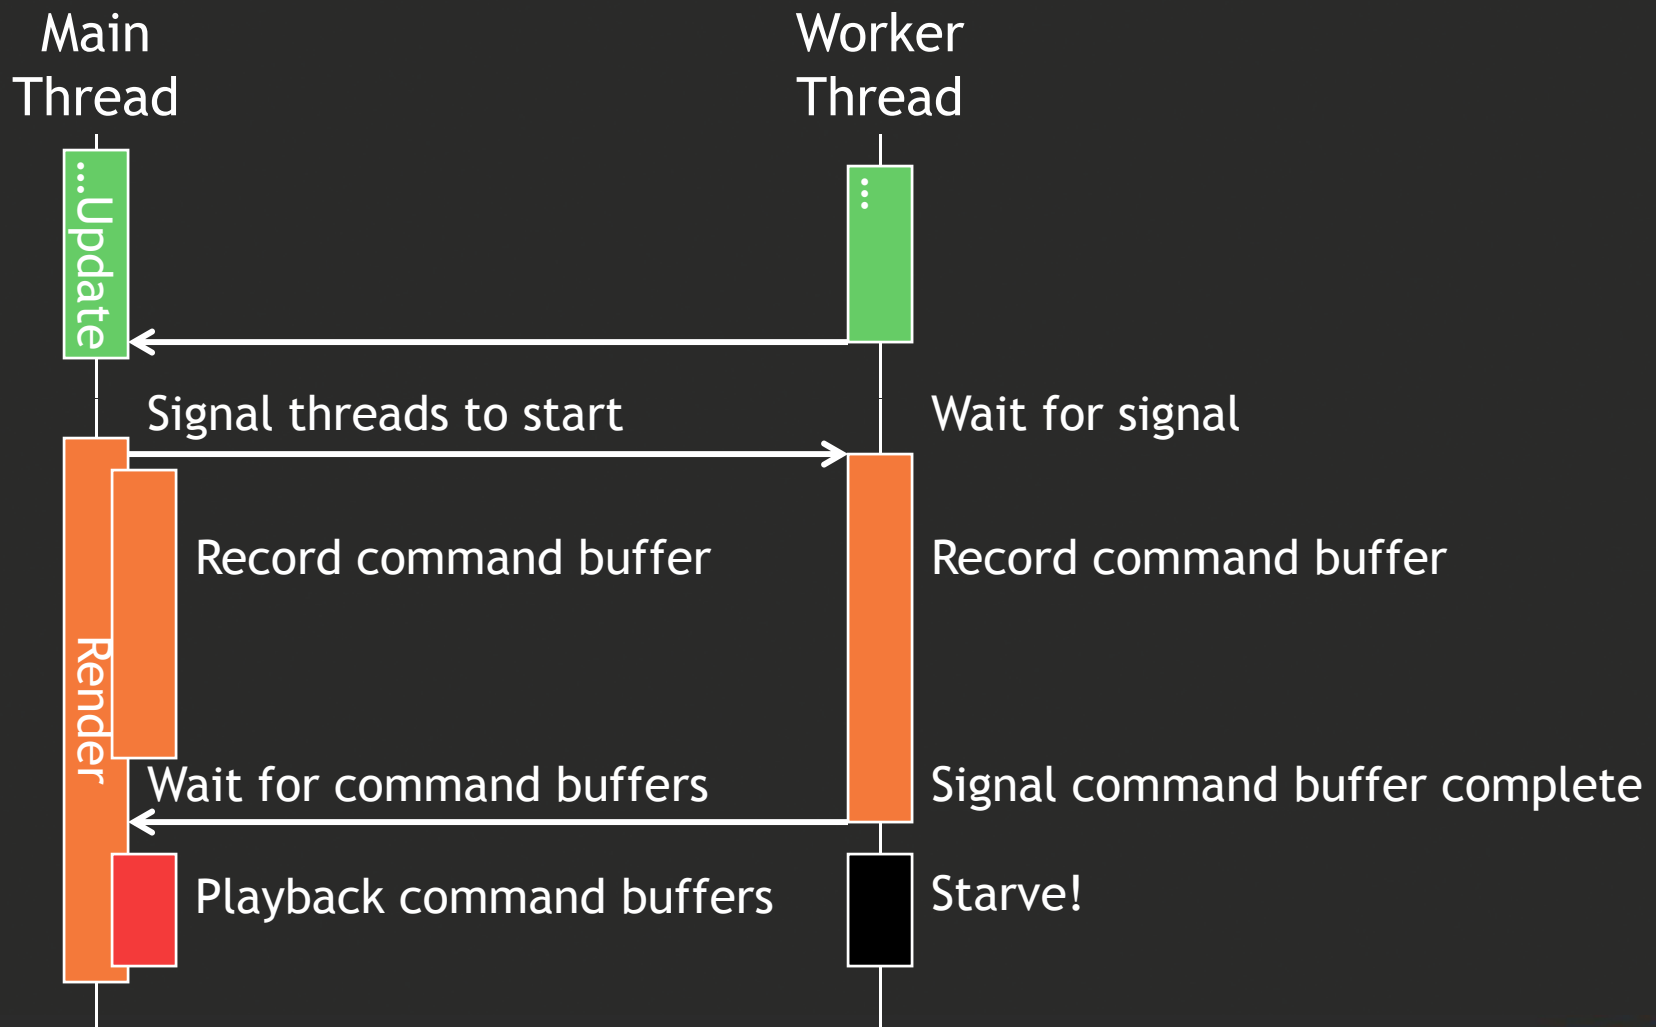


# Application Models

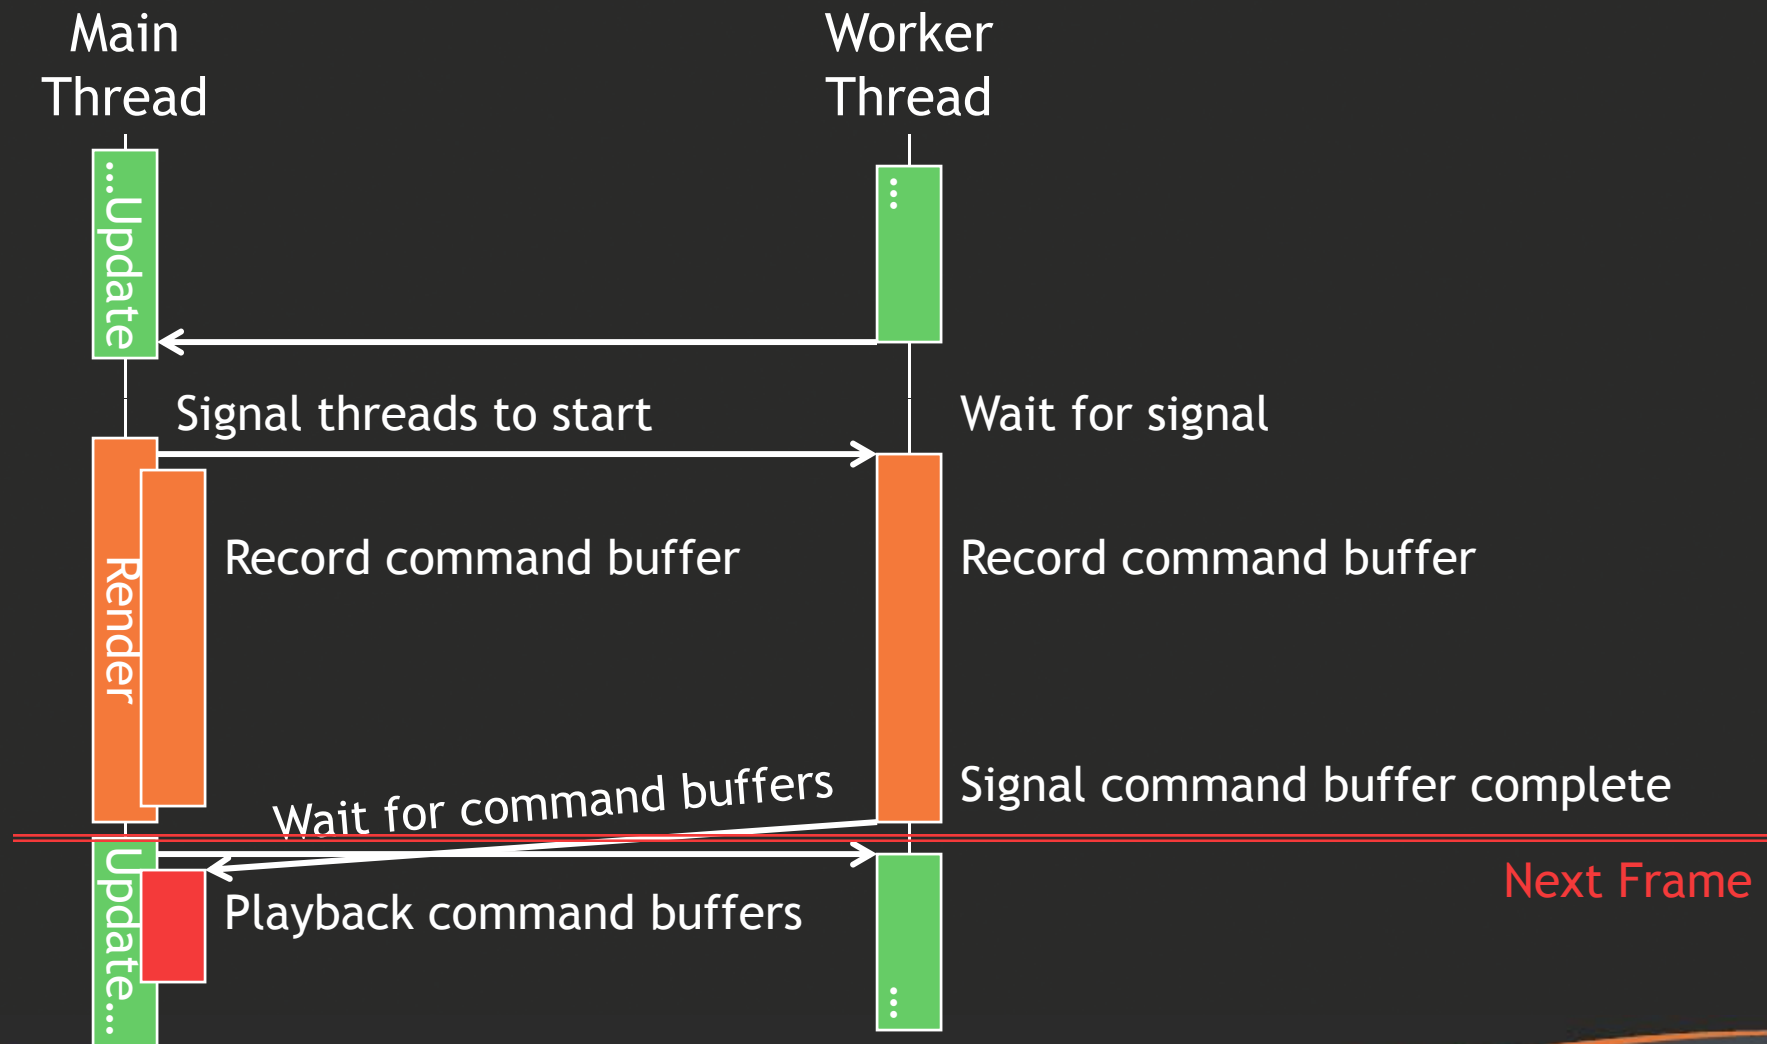
- Command buffers can be used in various ways by applications
  - Fork and join
  - Fork and join, frame deferred
  - Work queue
  - ...
- Record once, play back several times




# Fork & Join



# Fork & Join, Frame Deferred



# Work Queue



# Adapting to an Existing Codebase

- Refactor code to take pointer to device that can be changed easily
  - Easy if pointer passed on stack
  - Thread local storage if used from heap
- Add ownership of recording devices, playback class, and pool of command buffers
- Determine application model, and add high-level logic to parcel out rendering work.
- Manage resources over recording and playback

# Integration into DX Samples

- Instancing
  - Effects, shader constants
- Textures tutorial
  - Simple, added multithreading
- Stress test
  - Fork and join multithreading, with optional:
    - Frame delay of playback
    - Draw call count
    - CPU and memory access
    - Recorder thread count



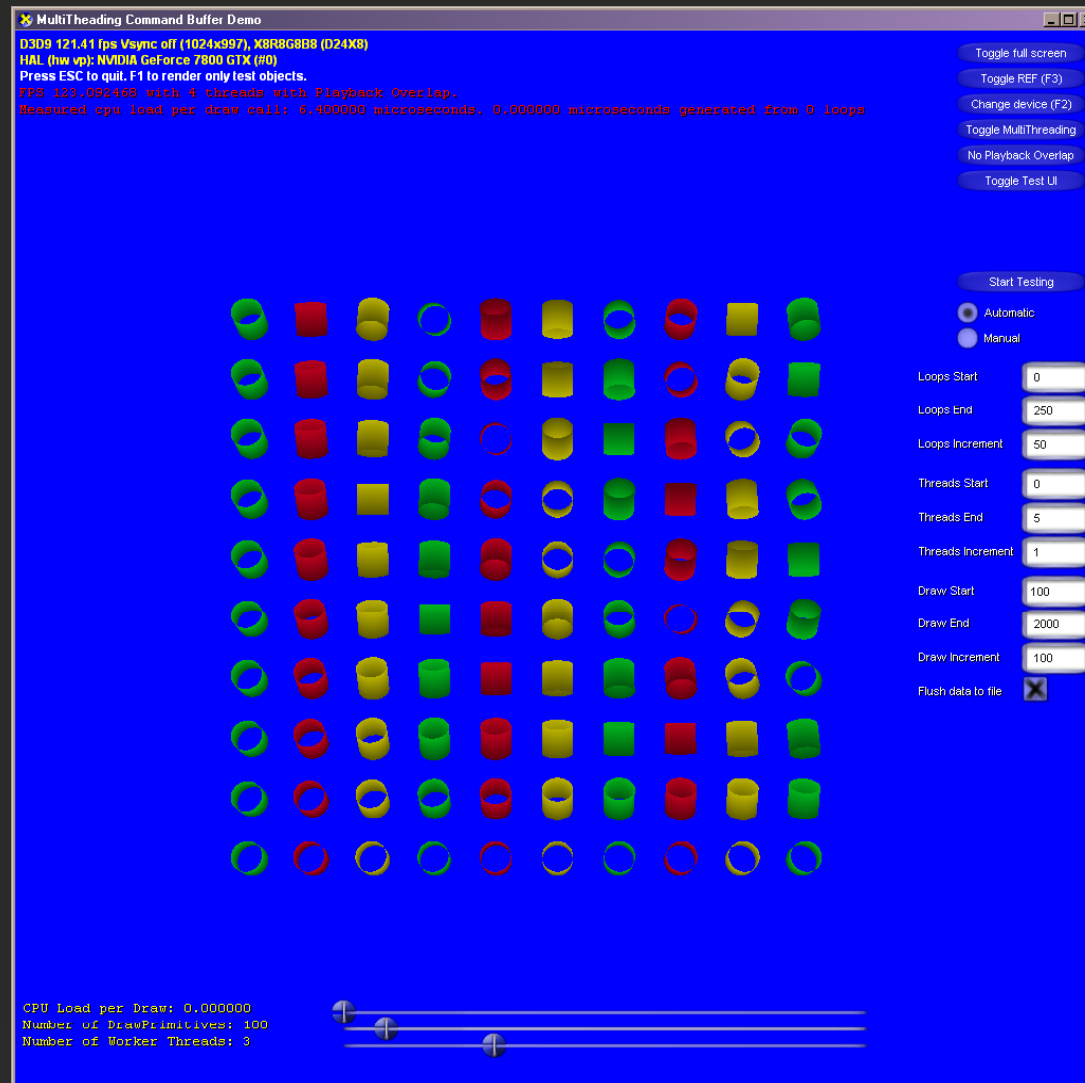
# Stress Test Information

- Render call contains:
  - Matrices computed with D3DX calls \* 3
  - SetTransform \* 3
  - SetRenderState
  - SetTexture
  - SetTextureStageState \* 8
  - SetStreamSource
  - SetFVF
  - DrawPrimitive

# CPU Busy Loops

- Draw call CPU cost varies in real applications
- Stress test simulates cost with CPU Busy Loops
  - Scattered reads from a large buffer in memory
  - Perform some logic, integer, and floating point operations
- Gamebryo render on DX9: 100-200  $\mu$ s
  - (on a Pentium 4, 3 GHz, nVidia 7800)
- Stress test can simulate Gamebryo render calls with 0-200 loops.

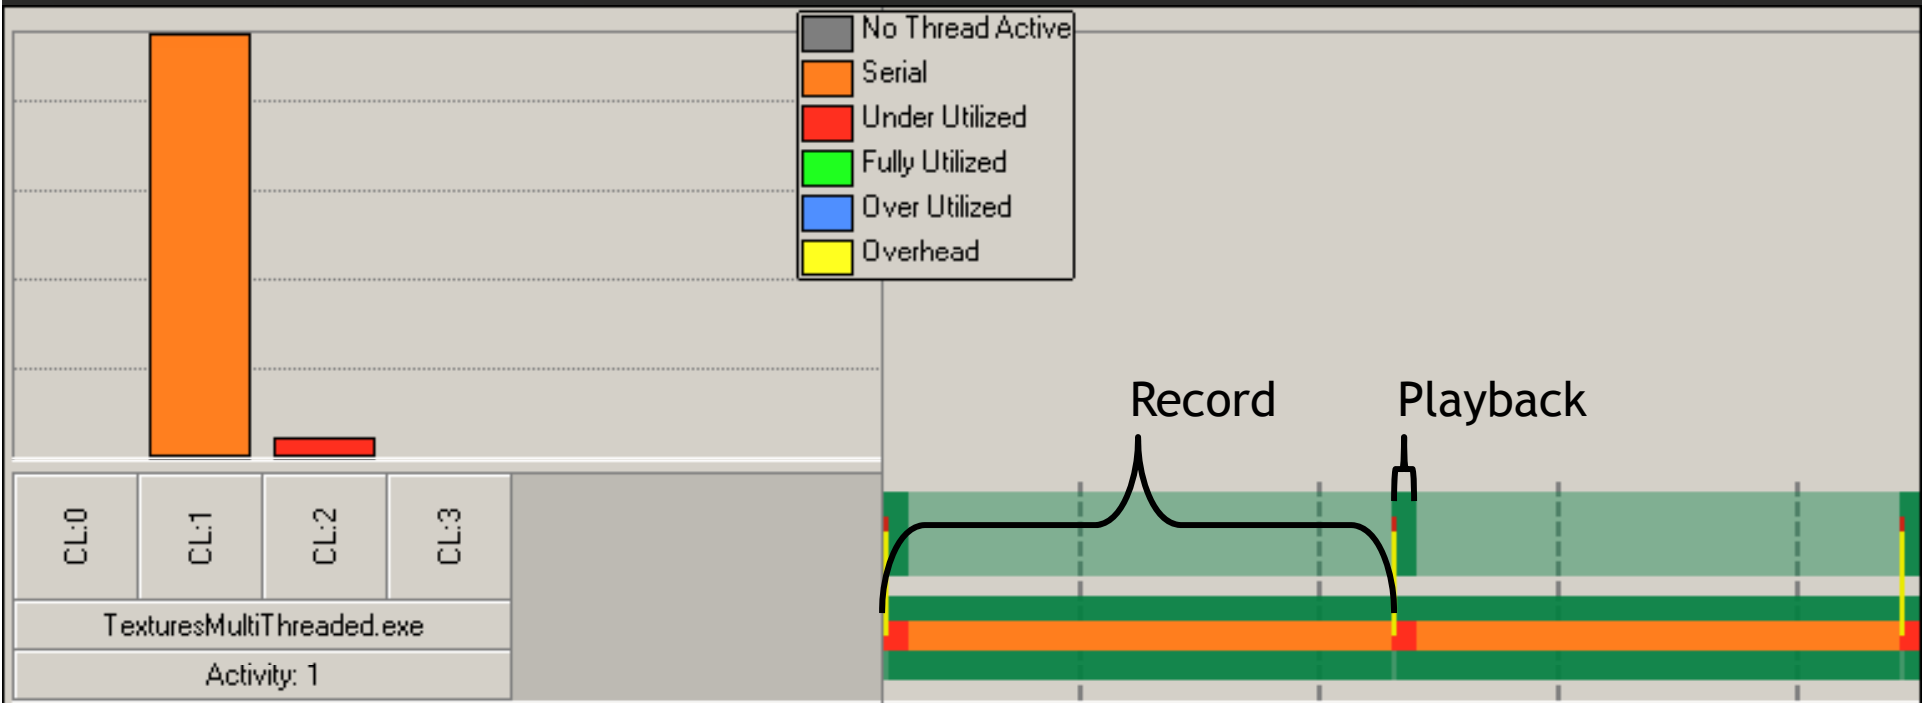
# DX Sample Stress Test Demo



# DX Call Cost vs. Recorder Cost

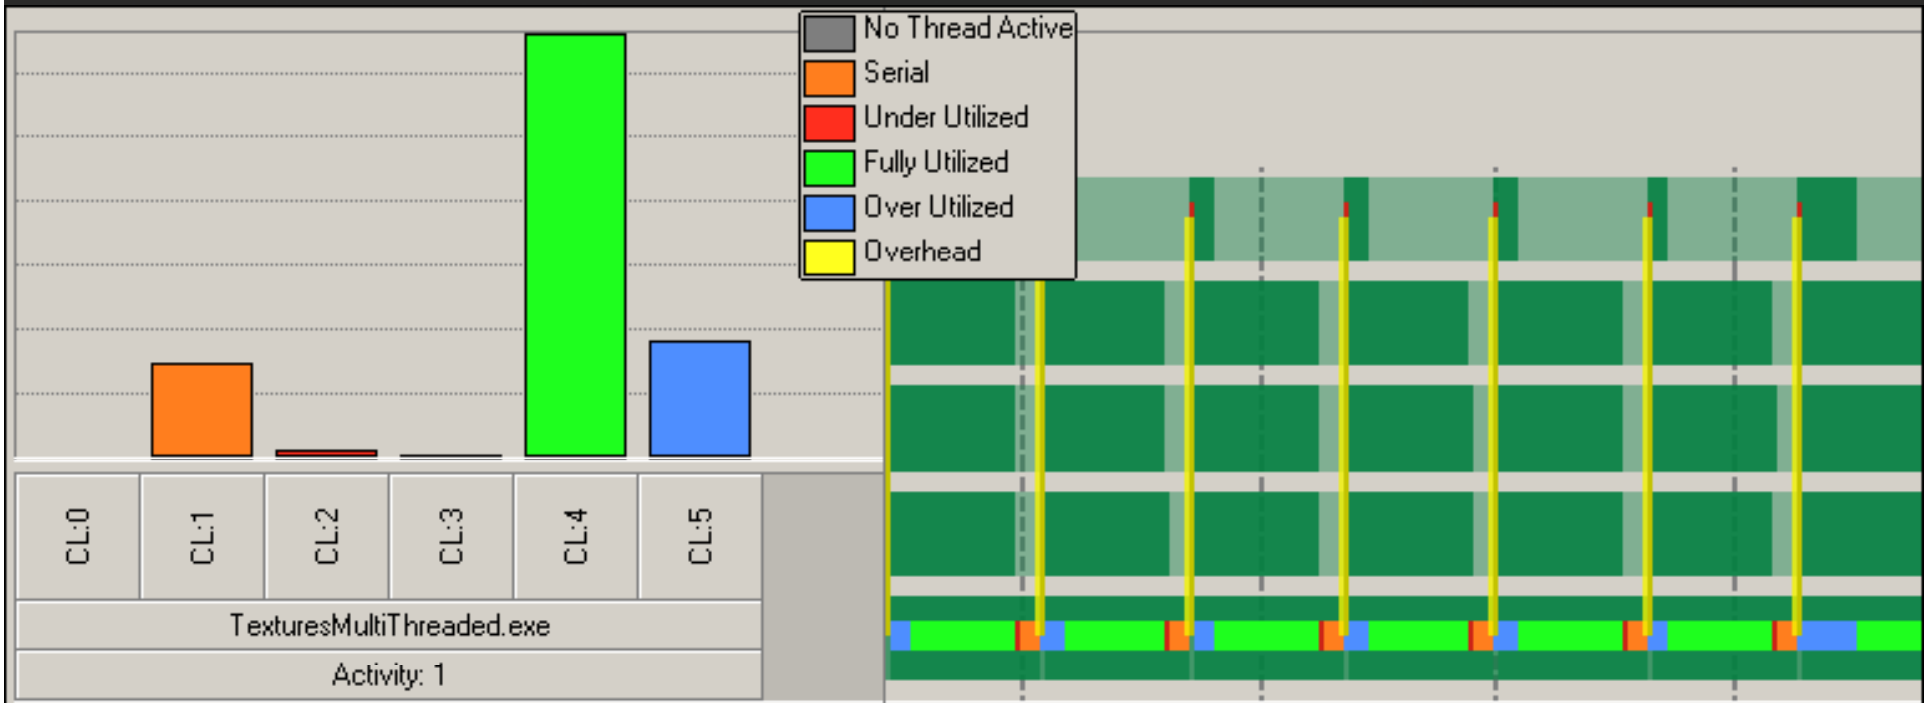
- Render call cost with DirectX device is 13 times as expensive as command buffer recorder
  - DX: 92 $\mu$ s
  - Recorder: 7 $\mu$ s
    - (on a Pentium 4, 3 GHz, nVidia 7800)

# Thread Profiler Quadcore 1 Recorder Thread

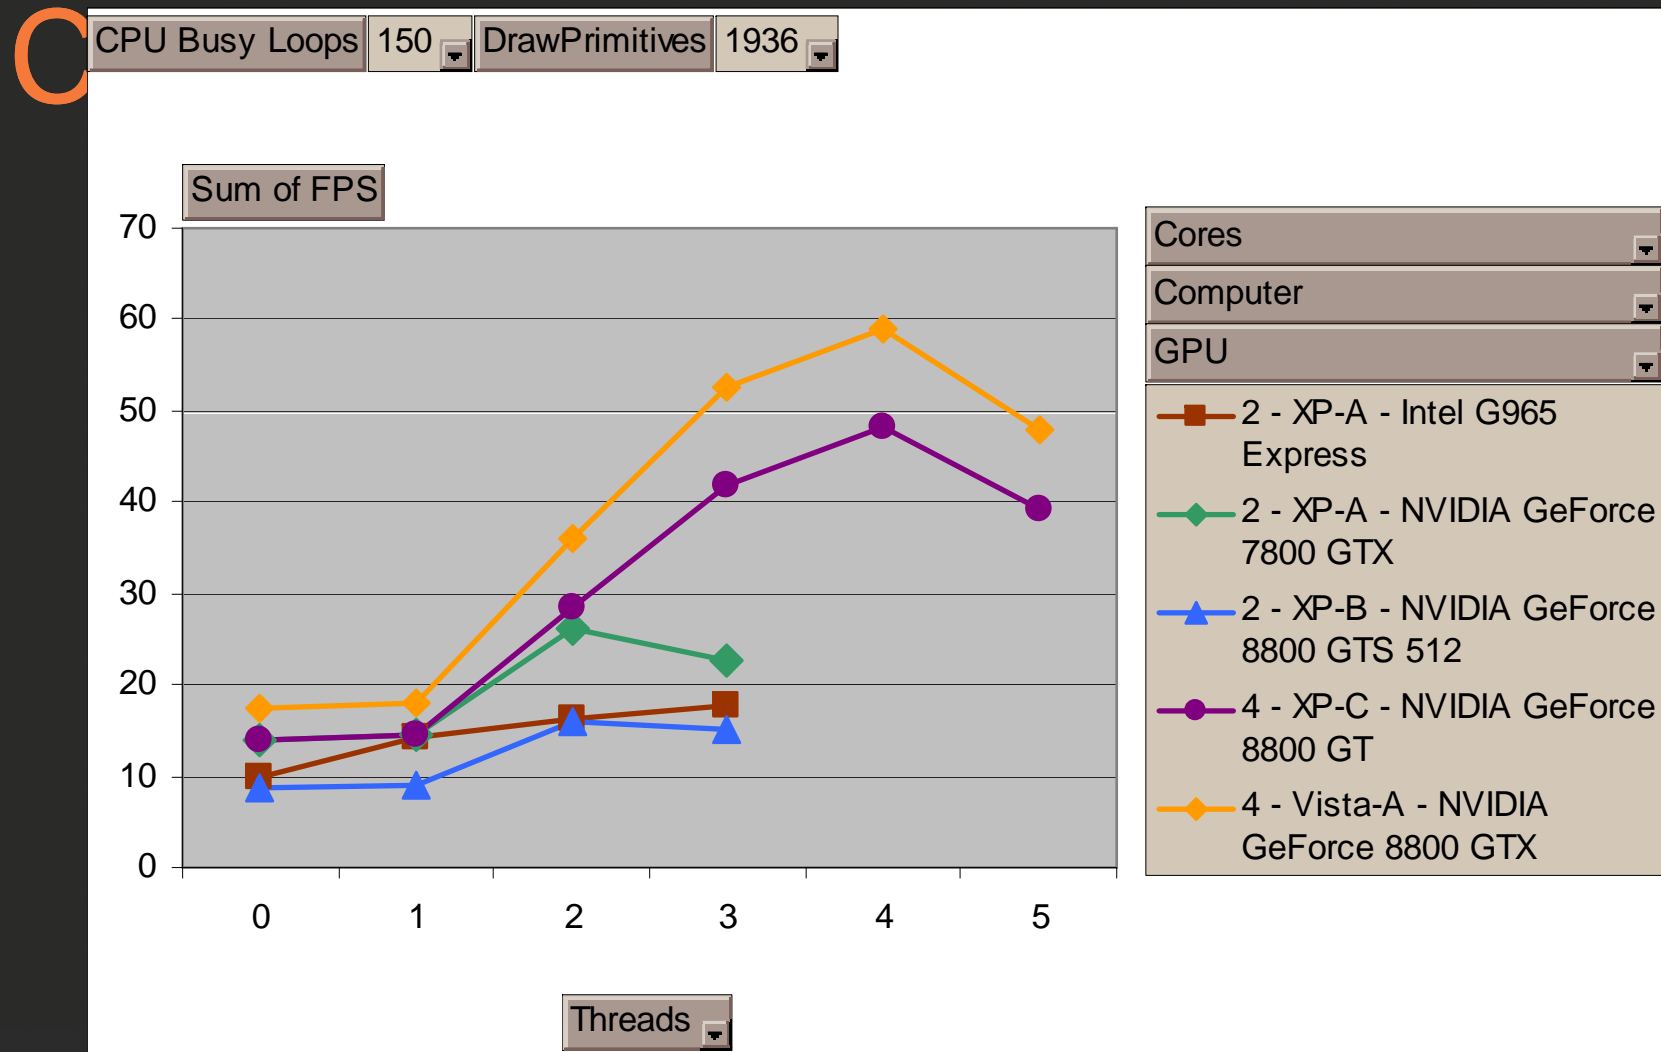




# Thread Profiler Quadcore 4 Recorder Threads




# FPS by Threads and




# Definition: Performance Ratio

- Charts that follow use  
Performance Ratio =  $\text{FPS}_{\text{test}} / \text{FPS}_{\text{baseline}}$
- Normalized result
- Useful for comparisons while varying
  - Number of draw calls
  - CPU busy loops

# Perf by Threads & Busy Loops




# Perf by Threads & Busy Loops

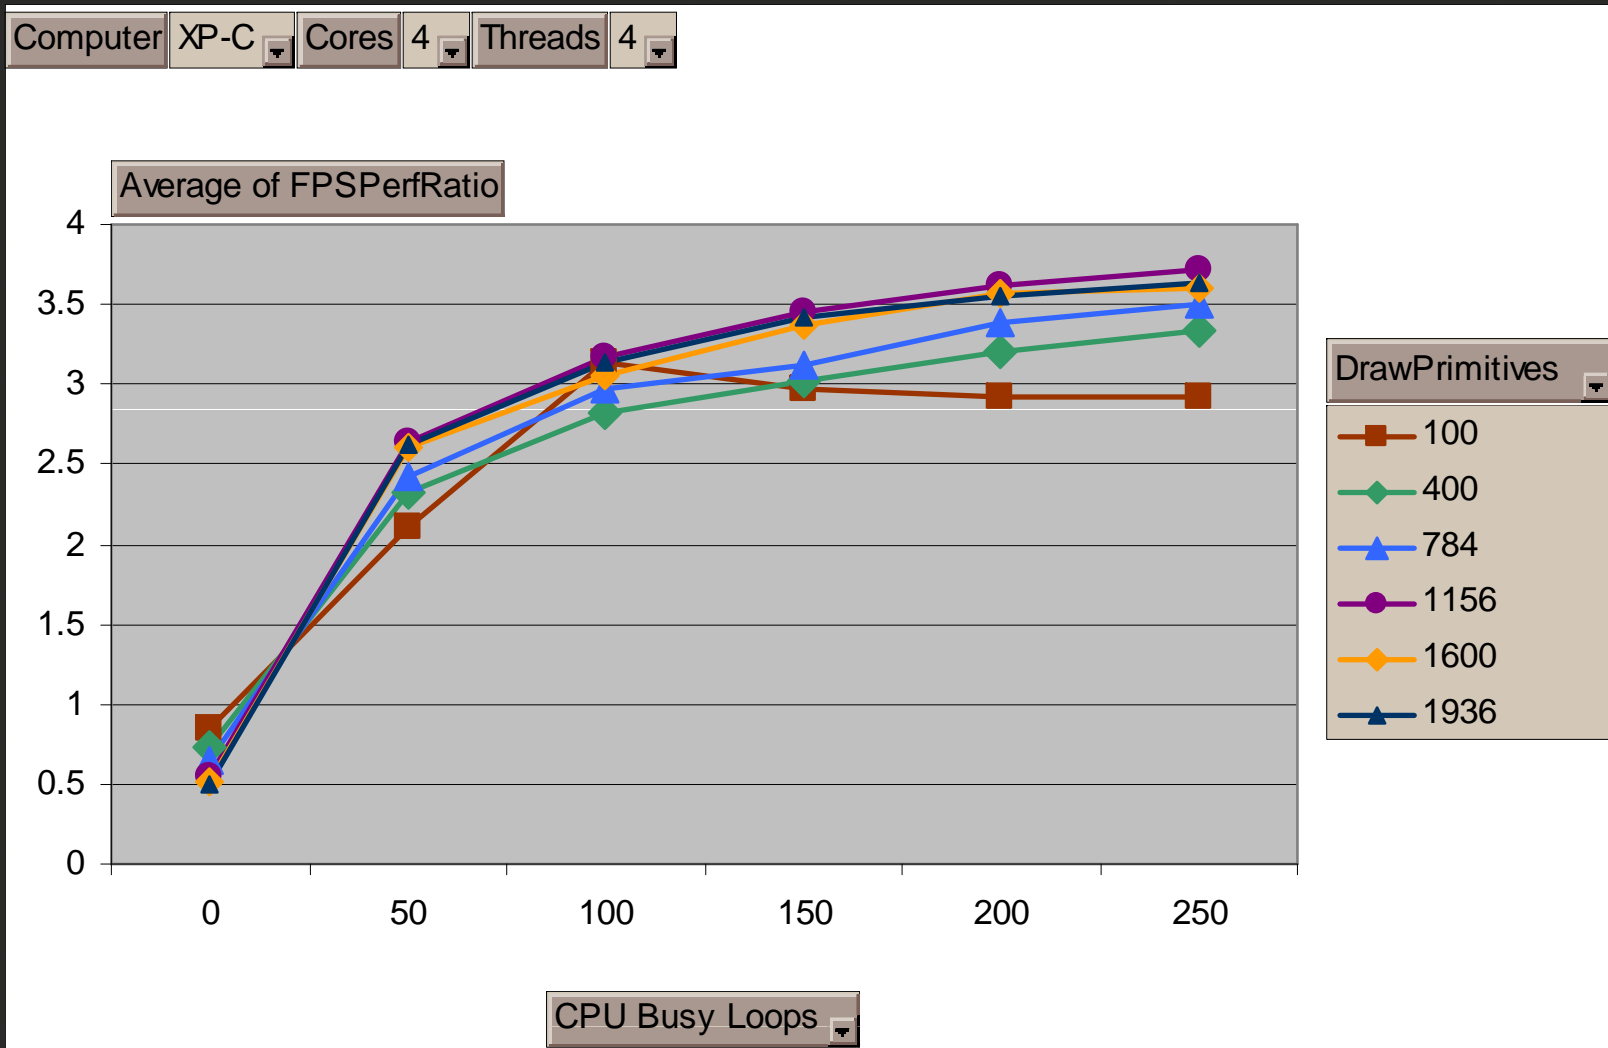




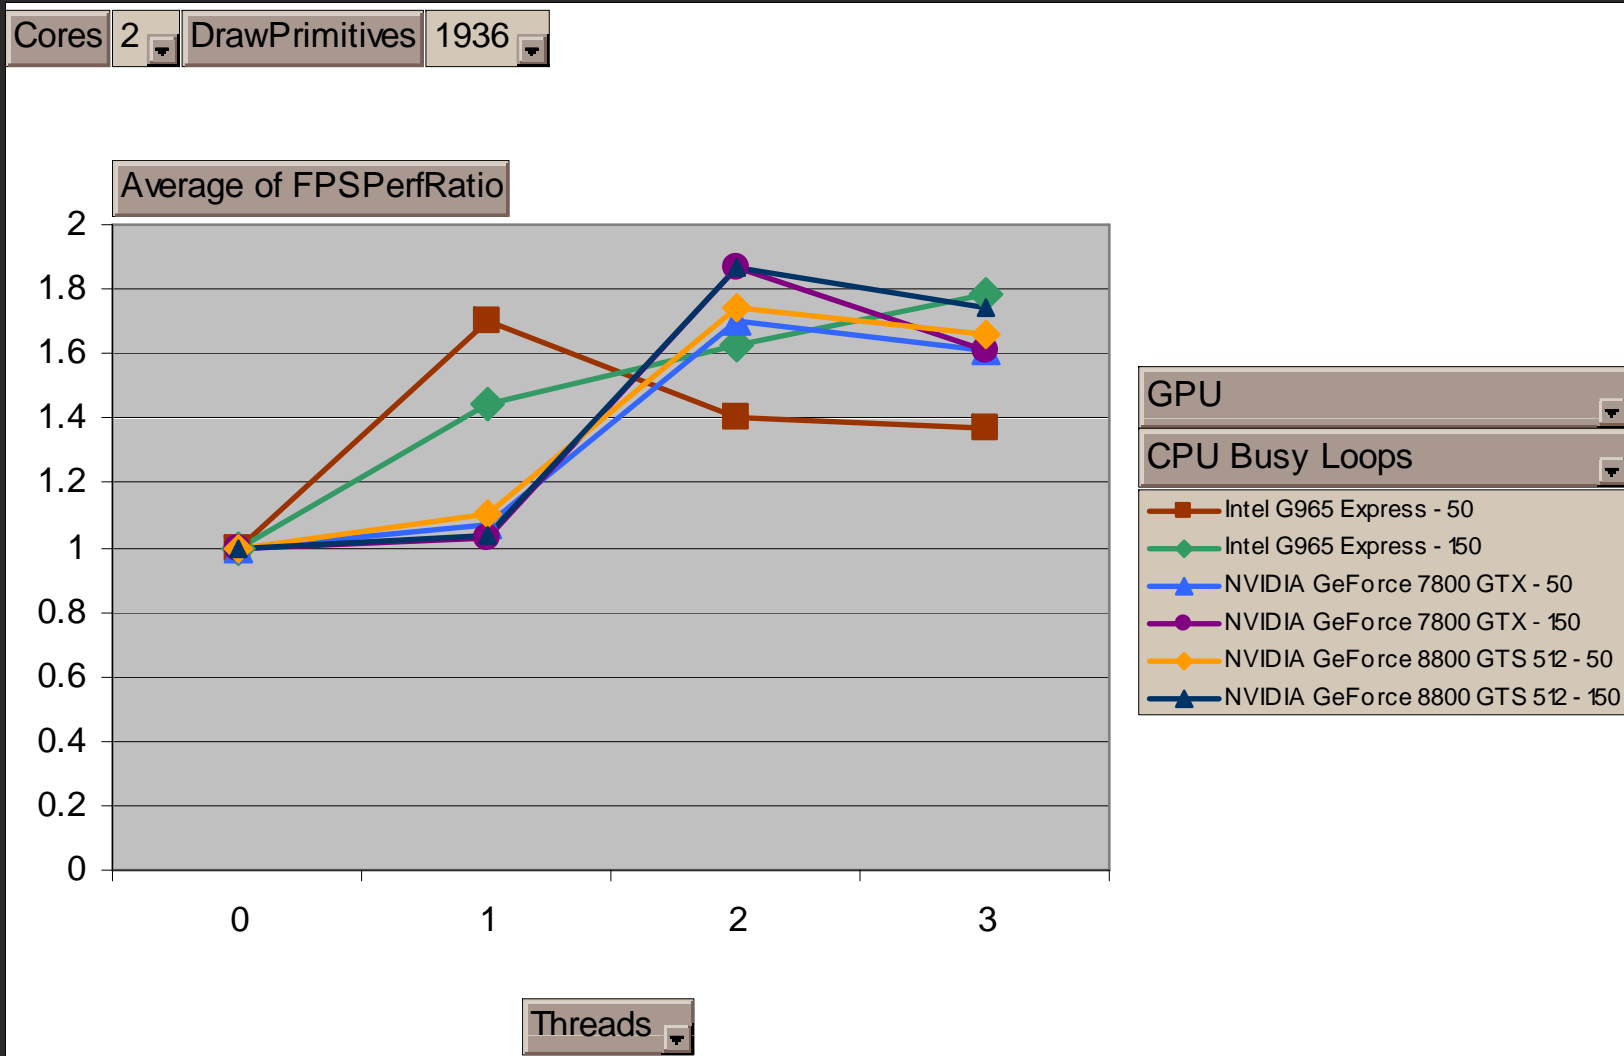
# Perf by Draws & Busy Loops



# Perf by Busy Loops & Draws



# Dual Core Results



# Future Work

- Resource management facilitated through command buffer, instead of application logic
- Optimization of command buffers by reordering order independent draw calls
- DirectX10

# Open Source Library

- Emergent has open sourced the command buffer library
  - Command buffer serialization
  - Recording device
  - Playback class
  - Redirecting device
  - EffectStateManager
  - DX9 only so far
- <http://emergent.net/GameFest2008>



# Thank You. Questions?

- Presented by:

- Vincent.Scheib@emergent.net

## Co-Developer:

- Bo.Wilson@emergent.net

- Consultation:

- Tim.Preston@emergent.net

- Dan.Amerson@emergent.net

- <http://emergent.net/GameFest2008>





# Gamefest

MICROSOFT GAME TECHNOLOGY CONFERENCE 2 0 0 8

[www.xnagamefest.com](http://www.xnagamefest.com)



© 2008 Microsoft Corporation. All rights reserved.  
This presentation is for informational purposes only.  
Microsoft makes no warranties, express or implied, in this summary.

Gamefest  
MICROSOFT GAME TECHNOLOGY CONFERENCE 2 0 0 8